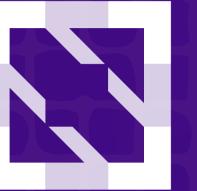




KubeCon



CloudNativeCon

 OPEN SOURCE SUMMIT

China 2019



KubeCon



CloudNativeCon



OPEN SOURCE SUMMIT

China 2019

Large Scale Distributed Deep Learning on Kubernetes Clusters

Yuan Tang – Ant Financial

Yong Tang – MobileIron

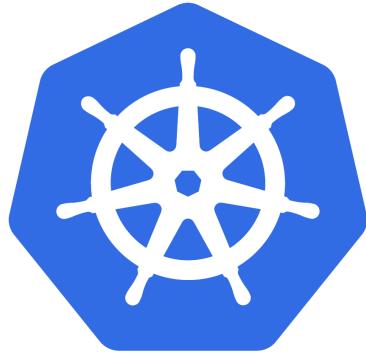


Speakers

- Yuan Tang
 - GitHub: terrytangyuan
 - Member: Kubeflow
 - Maintainer: MXNet, XGBoost, TensorFlow
 - Senior Software Engineer, Ant Financial
- Yong Tang
 - GitHub: yongtang
 - Maintainer: CoreDNS and Docker
 - Director of Engineering, MobileIron



TensorFlow

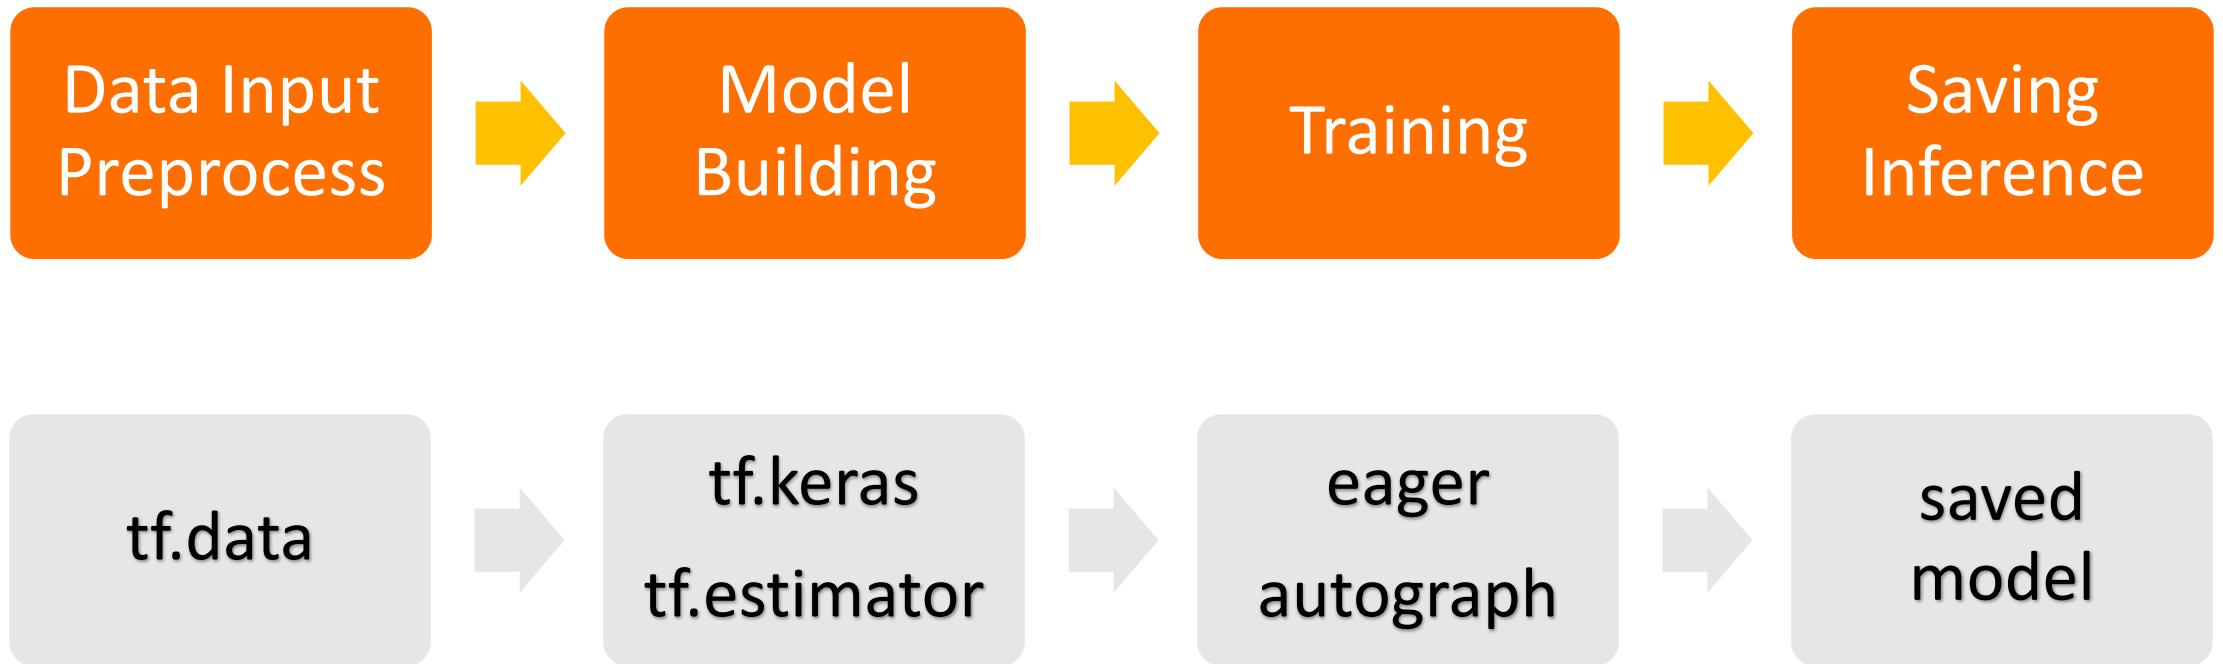


kubernetes

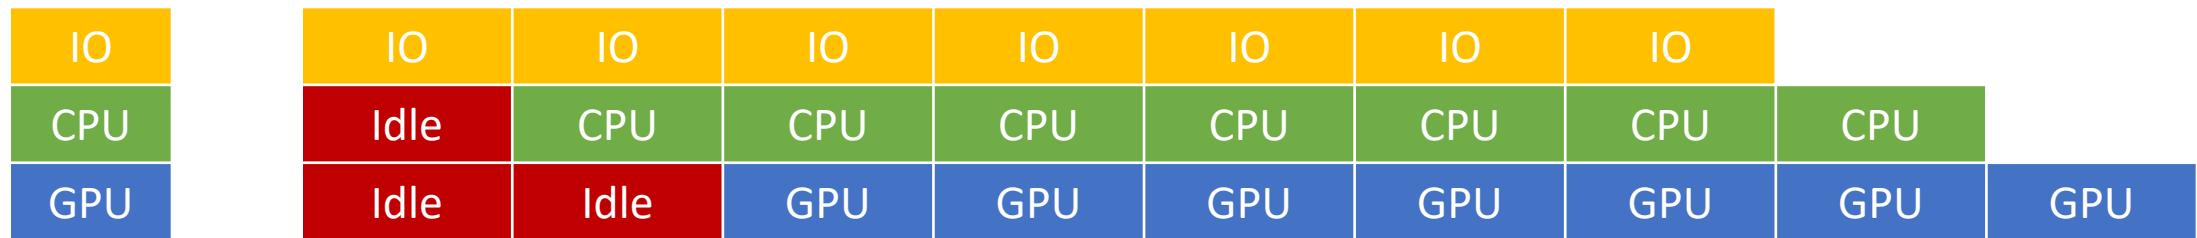
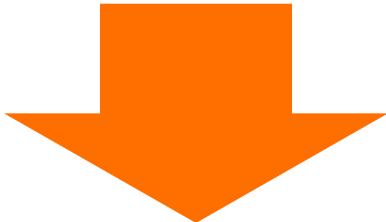
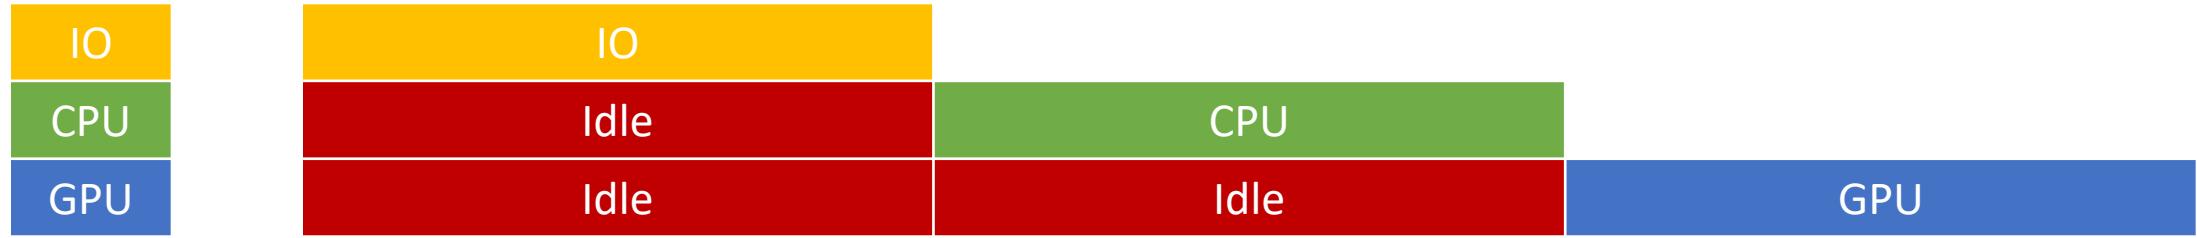


Kubeflow

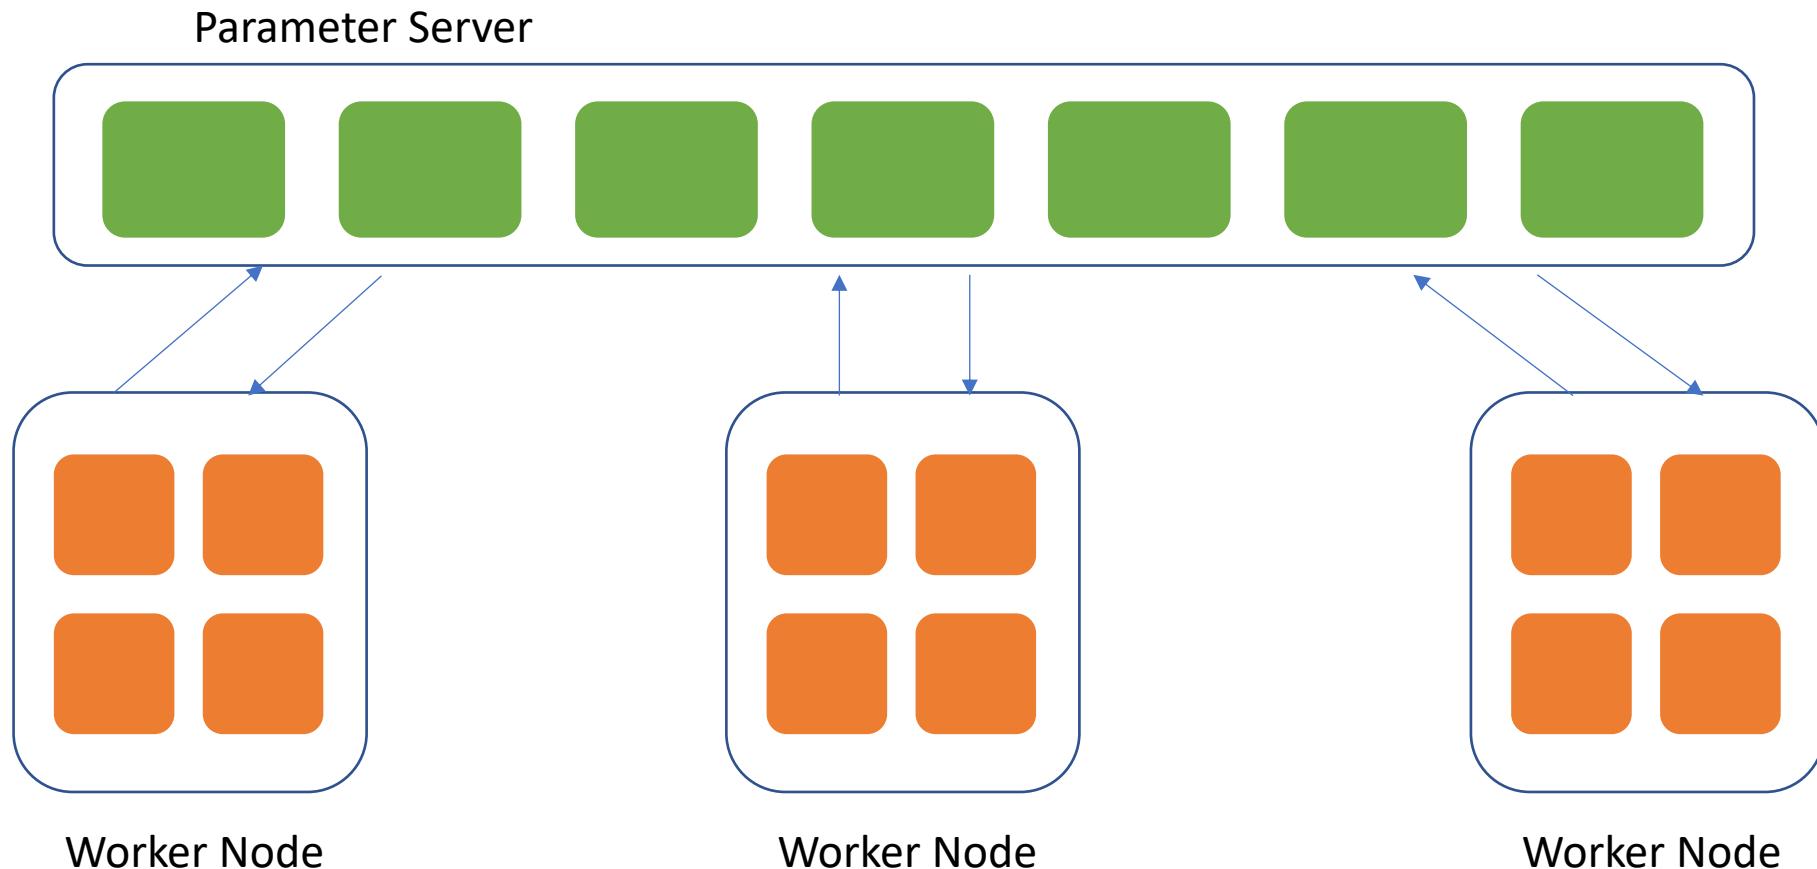
TensorFlow 2.0 Workflow



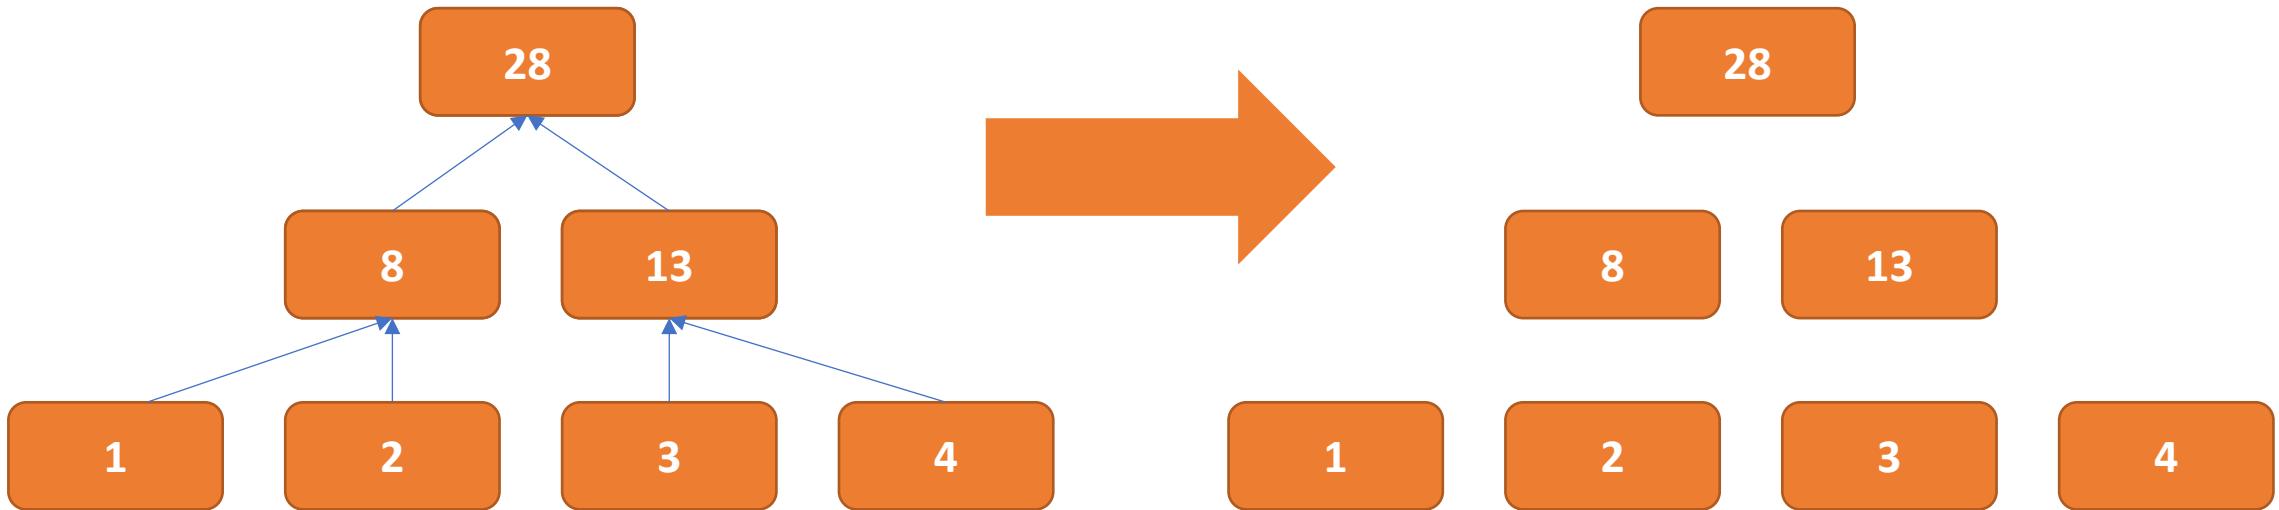
Orchestration for DL



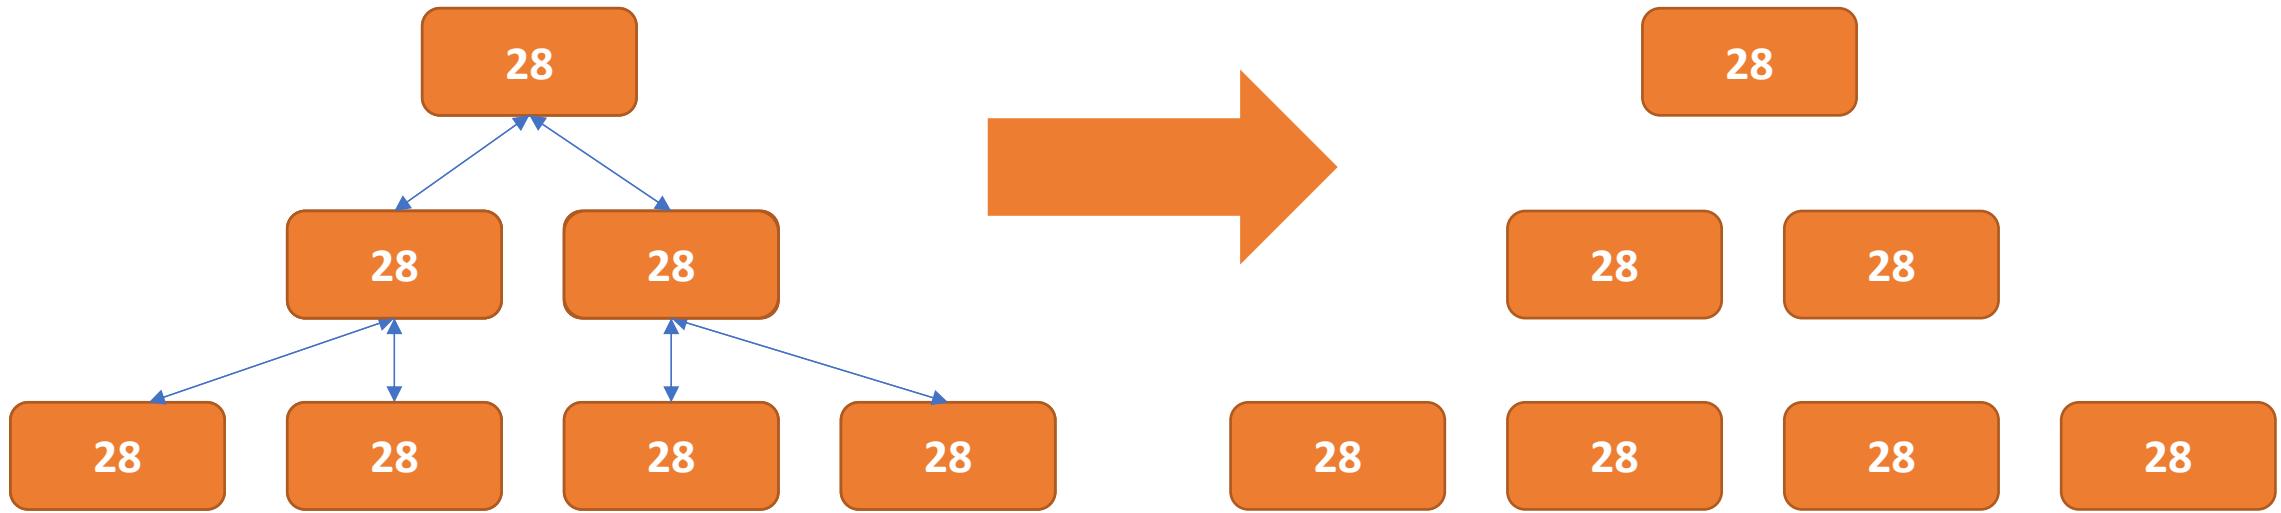
Parameter Server



Reduce



AllReduce



AllReduce == Reduce + Broadcast



Parameter Server

Parallelize on a machine

Parallelize in a cluster

Controversial

Cross device communication cost

Huge efforts invested over the years



Orchestration for DL



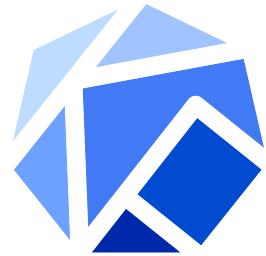
Stateful Metadata

Lifecycle Management

Kubernetes for Orchestration

Kubernetes Operators for ML

Kubernetes Operators



Kubeflow

TF Operator

PyTorch Operator

MPI Operator

Kubernetes Operators

	TF Operator	PyTorch Operator	MPI Operator
Framework Support	 TensorFlow	 PyTorch	 TF/Keras/MXNet/PyTorch OpenMPI
Distribution Strategy & Backend	<code>tf.distribute</code> MPI/NCCL/PS/TPU	<code>torch.distributed</code> Gloo/MPI/NCCL	<code>horovod</code> DistributedOptimizer (MPI Only)

TFJob vs. MPIJob

```
apiVersion: "kubeflow.org/v1beta1"
kind: TFJob
metadata:
  name: distributed-training
spec:
  tfReplicaSpecs:
    Worker:
      replicas: 4
      template:
        spec:
          containers:
            - name: tensorflow
              image: distributed_training_tf:latest
              resources:
                limits: nvidia.com/gpu: 4
              command: "python tf_benchmarks.py"
```

```
apiVersion: "kubeflow.org/v1alpha2"
kind: MPIJob
metadata:
  name: distributed-training
spec:
  mpiReplicaSpecs:
    Worker:
      replicas: 4
      template:
        spec:
          containers:
            - name: tensorflow
              image: distributed_training_hovorod:latest
              resources:
                limits: nvidia.com/gpu: 4
              command: "mpirun python hovorod_benchmarks.py"
```

TensorFlow

```
$ pip install tensorflow
$ pip install tensorflow-io
```

```
import tensorflow as tf
import tensorflow_io.mnist as mnist_io

dataset = mnist_io.MNISTDataset(image_filenames, label_filenames)
dataset = dataset.map(
    lambda x, y: (tf.image.convert_image_dtype(x, tf.float32), y)).batch(1000)

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(512, activation=tf.nn.relu),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(loss='mse', optimizer='sgd')
model.fit(dataset, epochs=2000)
model.evaluate(dataset)
```



Mirror Strategy in TF

```
import tensorflow as tf
import tensorflow_io.mnist as mnist_io

dataset = mnist_io.MNISTDataset(...)

model = tf.keras.Sequential([...])

mirrored_strategy = tf.distribute.MirroredStrategy()
with mirrored_strategy.scope():
    model.compile(loss='mse', optimizer='sgd')

model.fit(dataset, epochs=2000)
model.evaluate(dataset)
```

TensorFlow + Horovod

```
import tensorflow as tf
import tensorflow_io.mnist as mnist_io
import horovod.keras as hvd

dataset = mnist_io.MNISTDataset(...)

model = tf.keras.Sequential(...)

opt = tf.train.AdagradOptimizer(0.01 * hvd.size())
opt = hvd.DistributedOptimizer(opt)

model.compile(loss='mse', optimizer=opt)

callbacks = [
    hvd.callbacks.BroadcastGlobalVariablesCallback(0),
]
model.fit(dataset, epochs=2000, callbacks=callbacks)
model.evaluate(dataset)
```

TensorFlow vs. Horovod

```
import tensorflow as tf
import tensorflow_io.mnist as mnist_io

dataset = mnist_io.MNISTDataset(...)

model = tf.keras.Sequential([...])

mirrored_strategy = tf.distribute.MirroredStrategy()
with mirrored_strategy.scope():
    model.compile(loss='mse', optimizer='sgd')

model.fit(dataset, epochs=2000)
model.evaluate(dataset)
```

```
import tensorflow as tf
import tensorflow_io.mnist as mnist_io
import horovod.keras as hvd

dataset = mnist_io.MNISTDataset(...)

model = tf.keras.Sequential([...])

opt = tf.train.AdagradOptimizer(0.01 * hvd.size())
opt = hvd.DistributedOptimizer(opt)

model.compile(loss='mse', optimizer=opt)

callbacks = [
    hvd.callbacks.BroadcastGlobalVariablesCallback(0),
]
model.fit(dataset, epochs=2000, callbacks=callbacks)
model.evaluate(dataset)
```

PyTorch + Horovod

```
import torch
import horovod.torch as hvd

data_loader = torch.utils.data.DataLoader(train_dataset, batch_size=100)

model = ...

optimizer = torch.optim.SGD(model.parameters())
optimizer = hvd.DistributedOptimizer(
    optimizer, named_parameters=model.named_parameters())
hvd.broadcast_parameters(model.state_dict(), root_rank=0)

for epoch in range(100):
    for batch_idx, (data, target) in enumerate(data_loader):
        optimizer.zero_grad()
        output = model(data)
        loss = torch.nn.functional.F.nll_loss(output, target)
        loss.backward()
        optimizer.step()
```

Recall: TFJob vs. MPIJob

```
apiVersion: "kubeflow.org/v1beta1"
kind: TFJob
metadata:
  name: distributed-training
spec:
  tfReplicaSpecs:
    Worker:
      replicas: 4
      template:
        spec:
          containers:
            - name: tensorflow
              image: distributed_training_tf:latest
            resources:
              limits: nvidia.com/gpu: 4
            command: "python tf_benchmarks.py"
```

```
apiVersion: "kubeflow.org/v1alpha2"
kind: MPIJob
metadata:
  name: distributed-training
spec:
  mpiReplicaSpecs:
    Worker:
      replicas: 4
      template:
        spec:
          containers:
            - name: tensorflow
              image: distributed_training_hovorod:latest
            resources:
              limits: nvidia.com/gpu: 4
            command: "mpirun python hovorod_benchmarks.py"
```

Shared API and Best Practices



KubeCon

CloudNativeCon

OPEN SOURCE SUMMIT

China 2019

kubeflow / common

Unwatch 9 Unstar 11 Fork 8

Code Issues 6 Pull requests 2 Projects 0 Wiki Insights

Common APIs and libraries shared by other Kubeflow operator repositories.

39 commits 1 branch 0 releases 5 contributors Apache-2.0

Branch: master New pull request Create new file Upload files Find File Clone or download

terrytangyuan and k8s-ci-robot Correct function names in the comment (#32) Latest commit f38f5dc 2 days ago

File	Description	Time Ago
client	Common job controller library (#5)	28 days ago
hack	Common job controller library (#5)	28 days ago
job_controller	Correct function names in the comment (#32)	2 days ago
operator/v1	Fix incorrect name for restart policy exit code (#20)	10 days ago
test_job/v1	Remove mentions of tensorflow in test job (#21)	10 days ago
test_util/v1	chore: Fix package name (#27)	8 days ago
util	Move public util functions to util/status.go	6 days ago
.gitignore	Added .gitignore file (#16)	15 days ago
.travis.yml	Update geveralls ignore pattern	9 days ago
LICENSE	Create LICENSE	a month ago
OWNERS	Add terrytangyuan to OWNERS	a month ago
README.md	Add Travis badge and Go report card (#9)	27 days ago



Shared API and Best Practices



KubeCon



CloudNativeCon



OPEN SOURCE SUMMIT

China 2019

Common and standardized API spec

Base JobController interface

JobController implementation utilities

Testing utilities

THANK YOU